

Assignment_3_1

June 12, 2020

1 Assignment 3.1. Sequence Classification

2 Task: Aspect-level Sentiment Classification(10pt)

Group 82 * Student 1 : Joris Willems + 0908753 * Student 2 : Lars Schilders + 0908729

Reading material: - [1] R. He, WS. Lee & D. Dahlmeier. Exploiting document knowledge for aspect-level sentiment classification. 2018. <https://arxiv.org/abs/1806.04346>.

Build an attention-based aspect-level sentiment classification model with biLSTM. Your model shall include:

- BiLSTM network that learns sentence representation from input sequences.
- Attention network that assigns attention score over a sequence of biLSTM hidden states based on aspect terms representation.
- Fully connected network that predicts sentiment label, given the representation weighted by the attention score.

Requirements:

- You shall train your model based on transferring learning. That is, you need first train your model on document-level examples. Then the learned weights will be used to initialize aspect-level model and fine tune it on aspect-level examples.
- You shall use the alignment score function in attention network as following expression:

$$f_{score}(h, t) = \tanh(h^T W_a t)$$

- You shall evaluate the trained model on the provided test set and show the accuracy on test set.

```
[1]: from google.colab import drive
drive.mount('/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Aawg%3Aoauth%3A2.0%3Aoob&response_type=code&scope=email%20https%3A%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3A%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3A%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3A%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly

Enter your authorization code:

```
.....  
Mounted at /drive
```

```
[1]: import os  
import sys  
import codecs  
import operator  
import numpy as np  
import re  
from time import time
```

```
[2]: import _pickle as cPickle
```

3 Load Data

```
[3]: def read_pickle(data_path, file_name):  
  
    f = open(os.path.join(data_path, file_name), 'rb')  
    read_file = cPickle.load(f)  
    f.close()  
  
    return read_file  
  
def save_pickle(data_path, file_name, data):  
  
    f = open(os.path.join(data_path, file_name), 'wb')  
    cPickle.dump(data, f)  
    print(" file saved to: %s"%(os.path.join(data_path, file_name)))  
    f.close()
```

```
[4]: aspect_path = '/drive/My Drive/Colab Notebooks/2IMM10 - Deep Learning/  
↳Assignment 3/aspect-level' #for lars  
aspect_path = '/drive/My Drive/deeplearning/data/aspect-level/' #for joris  
aspect_path = '../courseFiles/Practicals/Practical 5/data/aspect-level/' #for_  
↳local jupyter notebook  
  
vocab = read_pickle(aspect_path, 'all_vocab.pkl')  
  
train_x = read_pickle(aspect_path, 'train_x.pkl')  
train_y = read_pickle(aspect_path, 'train_y.pkl')  
dev_x = read_pickle(aspect_path, 'dev_x.pkl')  
dev_y = read_pickle(aspect_path, 'dev_y.pkl')  
test_x = read_pickle(aspect_path, 'test_x.pkl')  
test_y = read_pickle(aspect_path, 'test_y.pkl')  
  
train_aspect = read_pickle(aspect_path, 'train_aspect.pkl')
```

```

dev_aspect = read_pickle(aspect_path, 'dev_aspect.pkl')
test_aspect = read_pickle(aspect_path, 'test_aspect.pkl')

pretrain_data = read_pickle(aspect_path, 'pretrain_data.pkl')
pretrain_label = read_pickle(aspect_path, 'pretrain_label.pkl')

```

```

[5]: def DecodeSequence(idx):
    review = str([list(vocab.keys())[i] for i in train_x[idx] if i!=0])
    aspectTerm = [list(vocab.keys())[i] for i in train_aspect[idx]]
    target = ["Positive", "Negative", "Neutral"][int(np.
    ↳nonzero(train_y[idx])[0])]
    print("Full review: {}".format(review))
    print("Aspect terms: {}".format(aspectTerm))
    print("Target: {}".format(target))

[6]: class Dataiterator_doc():
    """
    1) Iteration over minibatches using next(); call reset() between epochs
    ↳to randomly shuffle the data
    2) Access to the entire dataset using all()
    """

    def __init__(self, X, y, seq_length=32, decoder_dim=300, batch_size=32):
        self.X = X
        self.y = y
        self.num_data = len(X) # total number of examples
        self.batch_size = batch_size # batch size
        self.reset() # initial: shuffling examples and set index to 0

    def __iter__(self): # iterates data
        return self

    def reset(self): # initials
        self.idx = 0
        self.order = np.random.permutation(self.num_data) # shuffling examples
        ↳by providing randomized ids

    def __next__(self): # return model inputs - outputs per batch
        X_ids = [] # hold ids per batch
        while len(X_ids) < self.batch_size:
            X_id = self.order[self.idx] # copy random id from initial shuffling
            X_ids.append(X_id)
            self.idx += 1 #
            if self.idx >= self.num_data: # exception if all examples of data
        ↳have been seen (iterated)

```

```

        self.reset()
        raise StopIteration()

    batch_X = self.X[np.array(X_ids)] # X values (encoder input) per batch
    batch_y = self.y[np.array(X_ids)] # y_in values (decoder input) per
    ↪ batch
    return batch_X, batch_y

    def all(self): # return all data examples
        return self.X, self.y
class Dataiterator_aspect():
    '''
        1) Iteration over minibatches using next(); call reset() between epochs
    ↪ to randomly shuffle the data
        2) Access to the entire dataset using all()
    '''

    def __init__(self, aspect_data, seq_length=32, decoder_dim=300,
    ↪ batch_size=32):

        len_aspect_data = len(aspect_data[0])
        #self.len_doc_data = len(doc_data[0])

        self.X_aspect = aspect_data[0]
        self.y_aspect = aspect_data[1]
        self.aspect_terms = aspect_data[2]
        self.num_data = len_aspect_data
        self.batch_size = batch_size # batch size
        self.reset() # initial: shuffling examples and set index to 0

    def __iter__(self): # iterates data
        return self

    def reset(self): # initials
        self.idx = 0
        self.order = np.random.permutation(self.num_data) # shuffling examples
    ↪ by providing randomized ids

    def __next__(self): # return model inputs - outputs per batch

        X_ids = [] # hold ids per batch
        while len(X_ids) < self.batch_size:
            X_id = self.order[self.idx] # copy random id from initial shuffling
            X_ids.append(X_id)
            self.idx += 1 #

```

```

        if self.idx >= self.num_data: # exception if all examples of data
            have been seen (iterated)
            self.reset()
            raise StopIteration()

        batch_X_aspect = self.X_aspect[np.array(X_ids)] # X values (encoder
            input) per batch
        batch_y_aspect = self.y_aspect[np.array(X_ids)] # y_in values (decoder
            input) per batch
        batch_aspect_terms = self.aspect_terms[np.array(X_ids)]

        return batch_X_aspect, batch_y_aspect, batch_aspect_terms

    def all(self): # return all data examples
        return self.X_aspect, self.y_aspect, self.aspect_terms

```

```

[7]: from tensorflow import keras
      from keras.models import Model
      from keras.layers import Input, Embedding, Dense, Lambda, Dropout,
          LSTM, Bidirectional
      from keras.layers import Reshape, Activation, RepeatVector, concatenate,
          Concatenate, Dot, Multiply, Add
      import keras.backend as K
      from keras.engine.topology import Layer
      from keras import initializers
      from keras import regularizers
      from keras import constraints

```

Using TensorFlow backend.

```

[8]: overal_maxlen = 82
      overal_maxlen_aspect = 7

```

4 Define Attention Network Layer

- Define class for Attention Layer
- You need to finish the code for calculating the attention weights

```

[9]: class Attention(Layer):
      def __init__(self, **kwargs):
          """
          Keras Layer that implements an Content Attention mechanism.
          Supports Masking.
          """

```

```

self.supports_masking = True
self.init = initializers.get('glorot_uniform')

super(Attention, self).__init__(**kwargs)

def build(self, input_shape):
    assert type(input_shape) == list

    self.steps = input_shape[0][1]

    self.W = self.add_weight(shape=(input_shape[0][-1], input_shape[1][-1]),
                              initializer=self.init,
                              name='{}_W'.format(self.name),)

    self.built = True

def compute_mask(self, input_tensor, mask=None):
    assert type(input_tensor) == list
    assert type(mask) == list
    return None

def call(self, input_tensor, mask=None):

    # output of BiLSTM for sentence (h in paper)
    x = input_tensor[0]  #(None, 82, 600)

    # output of word embedding layer of aspect terms (t in paper)
    aspect = input_tensor[1]  #(None, 300)

    # used to remove influence of padded value
    #mask = mask[0]

     #(None, 600)
    aspect = K.transpose(K.dot(self.W, K.transpose(aspect)))

     #(None, 1, 600)
    aspect = K.expand_dims(aspect, axis=-2)

     #(None, 82, 600)
    aspect = K.repeat_elements(aspect, self.steps, axis=1)

     #(None, 82)
    beta_vec = K.tanh(K.sum(x*aspect, axis=-1))

     #(None, 82)
    alpha = K.exp(beta_vec)

```

```

        alpha /= K.cast(K.sum(alpha, axis=1, keepdims=True) + K.epsilon(), K.
↪floatx())

    return alpha

def compute_output_shape(self, input_shape):
    print((input_shape[0][0], input_shape[0][1]))
    return (input_shape[0][0], input_shape[0][1]) #(None, 82)

##### van paper #####
def call2(self, input_tensor, mask=None):
    x = input_tensor[0]
    aspect = input_tensor[1]
    mask = mask[0]

    aspect = K.transpose(K.dot(self.W, K.transpose(aspect)))
    aspect = K.expand_dims(aspect, axis=-2)
    aspect = K.repeat_elements(aspect, self.steps, axis=1)
    eij = K.sum(x*aspect, axis=-1)

    if self.bias:
        b = K.repeat_elements(self.b, self.steps, axis=0)
        eij += b

    eij = K.tanh(eij)

    a = K.exp(eij)

    if mask is not None:
        a *= K.cast(mask, K.floatx())

    a /= K.cast(K.sum(a, axis=1, keepdims=True) + K.epsilon(), K.floatx())

    return a

```

```

[10]: class Average(Layer):

    def __init__(self, mask_zero=True, **kwargs):
        self.mask_zero = mask_zero
        self.supports_masking = True
        super(Average, self).__init__(**kwargs)

    def call(self, x, mask=None):
        if self.mask_zero:
            mask = K.cast(mask, K.floatx())

```

```

        mask = K.expand_dims(mask)
        x = x * mask
        return K.sum(x, axis=1) / (K.sum(mask, axis=1) + K.epsilon())
    else:
        return K.mean(x, axis=1)

    def compute_output_shape(self, input_shape):
        return (input_shape[0], input_shape[-1])

    def compute_mask(self, x, mask):
        return None

```

5 Establish computation Graph for model

- Input tensors
- Shared WordEmbedding layer
- Attention network layer
- Shared BiLSTM layer
- Shared fully connected layer(prediction layer)

```

[11]: dropout = 0.5
      recurrent_dropout = 0.1
      vocab_size = len(vocab)
      num_outputs = 3 # labels

```

5.1 Input tensors

```

[14]: #YOUR CODE HERE ##### Inputs #####
      sentence_input = Input(shape=(overall_maxlen,), dtype='int32',
      ↪name='sentence_input')
      aspect_input = Input(shape=(overall_maxlen_aspect,), dtype='int32',
      ↪name='aspect_input')
      pretrain_input = Input(shape=(None,), dtype='int32', name='pretrain_input')

```

5.2 Shared WordEmbedding layer

```

[15]: #YOUR CODE HERE### represent aspect as averaged word embedding ###
      word_emb = Embedding(vocab_size, 300, mask_zero=True, name='word_emb')
      aspect_term_embs = word_emb(aspect_input)
      aspect_embs = Average(mask_zero=True, name='aspect_emb')(aspect_term_embs)

```

```

[16]: #YOUR CODE HERE ### sentence representation from embedding ###
      sentence_embs = word_emb(sentence_input) # from aspect-level domain
      pretrain_embs = word_emb(pretrain_input) # from document-level domain

```


5.3 Shared BiLSTM layer

```
[17]: #YOUR CODE HERE ### sentence representation from embedding ###
bilstm = Bidirectional(LSTM(300, return_sequences=True, dropout=dropout,
    ↪ recurrent_dropout=recurrent_dropout), name='BiLSTM')
pretrain_lstm = bilstm(pretrain_embs) #pretrain the bilstm (no attention
    ↪ involved) MUST BE AVERAGED
sentence_lstm = bilstm(sentence_embs) #finetune bilstm with attention mechanism
    ↪ to obtain aspect-level info
```

5.4 Attention Layer

```
[18]: ##YOUR CODE HERE
attention = Attention()([sentence_lstm, aspect_embs])
```

(None, 82)

5.5 Prediction Layer

```
[19]: z = Dot(axes=1)([sentence_lstm, attention])

predictionLayer = Dense(3, activation='softmax')

aspect_probs = predictionLayer(z)

pretrain_avg = Average(mask_zero=True)(pretrain_lstm)
pretrain_probs = predictionLayer(pretrain_avg)
```

6 Build Models for document-level and aspect-level data

- The two models shared the embedding, BiLSTM, Prediction Layer

```
[22]: ### YOUR CODE HERE
model1 = Model(inputs=[pretrain_input], outputs=[pretrain_probs])

model2 = Model(inputs=[sentence_input, aspect_input], outputs=[aspect_probs])

model2.summary()
```

Model: "model_2"

```
-----
-----
Layer (type)                Output Shape          Param #   Connected to
=====
=====
sentence_input (InputLayer)  (None, 82)            0
```

aspect_input (InputLayer)	(None, 7)	0	
word_emb (Embedding)	multiple	3000900	
aspect_input[0][0]			
sentence_input[0][0]			
BiLSTM (Bidirectional)	multiple	1442400	word_emb[1][0]
aspect_emb (Average)	(None, 300)	0	word_emb[0][0]
attention_1 (Attention)	(None, 82)	180000	BiLSTM[1][0]
aspect_emb[0][0]			
dot_1 (Dot)	(None, 600)	0	BiLSTM[1][0]
attention_1[0][0]			
dense_1 (Dense)	(None, 3)	1803	dot_1[0][0]
=====			
Total params: 4,625,103			
Trainable params: 4,625,103			
Non-trainable params: 0			

7 Train Model

- First Train model on document-level data.
- Then Train model on aspect-level data

7.1 Train on document-level data

```
[ ]: import keras.optimizers as opt
from sklearn.model_selection import train_test_split

# split the training dataset to be able to stop training when overfitting
↳ starts to kick in
pretrain_data_train, pretrain_data_test, pretrain_label_train,
↳ pretrain_label_test = train_test_split(pretrain_data,
```

```

        pretrain_label,

        test_size=0.05,

        stratify=pretrain_label)

optimizer=opt.RMSprop(lr=0.0005, rho=0.9, epsilon=1e-06, clipnorm=10,
    ↪clipvalue=0)
model1.compile(optimizer=optimizer, loss='categorical_crossentropy',
    ↪metrics=['categorical_accuracy'])
batch_size = 256

train_steps_epoch = len(pretrain_data_train)/batch_size
batch_train_iter_doc = Dataiterator_doc(pretrain_data_train,
    ↪pretrain_label_train, batch_size=batch_size)

test_steps_epoch = len(pretrain_data_test)/batch_size
batch_test_iter_doc = Dataiterator_doc(pretrain_data_test, pretrain_label_test,
    ↪batch_size=batch_size)

```

```

[24]: ###YOUR CODE HERE###

from keras.callbacks import EarlyStopping, ModelCheckpoint

def train_gen_doc():
    while True:
        for batch in batch_train_iter_doc:
            yield batch

def test_gen_doc():
    while True:
        for batch in batch_test_iter_doc:
            yield batch

history1 = model1.fit_generator(train_gen_doc(),
                                steps_per_epoch=train_steps_epoch,
                                validation_data=test_gen_doc(),
                                validation_steps=test_steps_epoch,
                                epochs=20,
                                ↪callbacks = [EarlyStopping(monitor='val_loss',
                                ↪patience=2)]
                                )

# save model
model1.save(aspect_path + "../model1.h5")

```

/usr/local/lib/python3.6/dist-

```
packages/tensorflow/python/framework/indexed_slices.py:434: UserWarning:
Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may
consume a large amount of memory.
```

```
"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "
```

```
Epoch 1/20
```

```
112/111 [=====] - 262s 2s/step - loss: 0.9200 -
categorical_accuracy: 0.5683 - val_loss: 0.7362 - val_categorical_accuracy:
0.6628
```

```
Epoch 2/20
```

```
112/111 [=====] - 263s 2s/step - loss: 0.7161 -
categorical_accuracy: 0.6853 - val_loss: 0.6919 - val_categorical_accuracy:
0.6914
```

```
Epoch 3/20
```

```
112/111 [=====] - 264s 2s/step - loss: 0.6462 -
categorical_accuracy: 0.7210 - val_loss: 0.7817 - val_categorical_accuracy:
0.6934
```

```
Epoch 4/20
```

```
112/111 [=====] - 263s 2s/step - loss: 0.6059 -
categorical_accuracy: 0.7405 - val_loss: 0.7327 - val_categorical_accuracy:
0.7031
```

7.2 Train on aspect-level data

```
[ ]: batch_size = 128

train_steps_epoch = len(train_x)/batch_size
batch_train_iter_aspect = Dataiterator_aspect([train_x, train_y, train_aspect],
↪batch_size=batch_size)
val_steps_epoch = len(dev_x)/batch_size
batch_val_iter_aspect = Dataiterator_aspect([dev_x, dev_y, dev_aspect],
↪batch_size=batch_size)

import keras.optimizers as opt
optimizer = opt.Adam(lr=0.00005, beta_1=0.9, beta_2=0.999, epsilon=1e-08,
↪clipnorm=10, clipvalue=0)
model2.compile(optimizer=optimizer, loss='categorical_crossentropy',
↪metrics=['categorical_accuracy'])
```

```
[26]: ### YOUR CODE HERE

def train_gen_aspect():
    while True:
        train_batches = [[X, aspect], [y]] for X, y, aspect in
↪batch_train_iter_aspect

        for batch in train_batches:
```

```

        yield batch

def val_gen_aspect():
    while True:
        val_batches = [[[X, aspect], [y]] for X, y, aspect in
↪batch_val_iter_aspect]

        for batch in val_batches:
            yield batch

history2 = model2.fit_generator(train_gen_aspect(),
                                steps_per_epoch=train_steps_epoch,
                                epochs=25,
                                validation_data=val_gen_aspect(),
                                validation_steps=val_steps_epoch
                                )

```

/usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/indexed_slices.py:434: UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Epoch 1/25

15/14 [=====] - 12s 798ms/step - loss: 0.9739 - categorical_accuracy: 0.5594 - val_loss: 1.0163 - val_categorical_accuracy: 0.5586

Epoch 2/25

15/14 [=====] - 10s 698ms/step - loss: 0.9114 - categorical_accuracy: 0.5833 - val_loss: 0.8177 - val_categorical_accuracy: 0.6172

Epoch 3/25

15/14 [=====] - 11s 703ms/step - loss: 0.8659 - categorical_accuracy: 0.6094 - val_loss: 0.8023 - val_categorical_accuracy: 0.5996

Epoch 4/25

15/14 [=====] - 10s 700ms/step - loss: 0.8493 - categorical_accuracy: 0.6182 - val_loss: 0.9249 - val_categorical_accuracy: 0.6211

Epoch 5/25

15/14 [=====] - 10s 696ms/step - loss: 0.8144 - categorical_accuracy: 0.6380 - val_loss: 0.8567 - val_categorical_accuracy: 0.6367

Epoch 6/25

15/14 [=====] - 10s 694ms/step - loss: 0.8110 - categorical_accuracy: 0.6484 - val_loss: 0.9310 - val_categorical_accuracy: 0.6562

Epoch 7/25

15/14 [=====] - 10s 694ms/step - loss: 0.7958 -
categorical_accuracy: 0.6505 - val_loss: 0.8238 - val_categorical_accuracy:
0.6406
Epoch 8/25
15/14 [=====] - 10s 697ms/step - loss: 0.7703 -
categorical_accuracy: 0.6635 - val_loss: 0.8402 - val_categorical_accuracy:
0.6348
Epoch 9/25
15/14 [=====] - 10s 700ms/step - loss: 0.7614 -
categorical_accuracy: 0.6687 - val_loss: 0.7398 - val_categorical_accuracy:
0.6543
Epoch 10/25
15/14 [=====] - 10s 687ms/step - loss: 0.7421 -
categorical_accuracy: 0.6885 - val_loss: 0.7482 - val_categorical_accuracy:
0.6699
Epoch 11/25
15/14 [=====] - 10s 695ms/step - loss: 0.7427 -
categorical_accuracy: 0.6703 - val_loss: 0.8124 - val_categorical_accuracy:
0.6387
Epoch 12/25
15/14 [=====] - 10s 691ms/step - loss: 0.7210 -
categorical_accuracy: 0.6984 - val_loss: 0.8509 - val_categorical_accuracy:
0.6777
Epoch 13/25
15/14 [=====] - 11s 720ms/step - loss: 0.7122 -
categorical_accuracy: 0.6938 - val_loss: 0.9618 - val_categorical_accuracy:
0.6582
Epoch 14/25
15/14 [=====] - 11s 700ms/step - loss: 0.6920 -
categorical_accuracy: 0.7115 - val_loss: 0.8218 - val_categorical_accuracy:
0.6738
Epoch 15/25
15/14 [=====] - 10s 695ms/step - loss: 0.6884 -
categorical_accuracy: 0.7141 - val_loss: 0.7906 - val_categorical_accuracy:
0.6777
Epoch 16/25
15/14 [=====] - 10s 699ms/step - loss: 0.6769 -
categorical_accuracy: 0.7271 - val_loss: 0.7560 - val_categorical_accuracy:
0.6738
Epoch 17/25
15/14 [=====] - 10s 692ms/step - loss: 0.6476 -
categorical_accuracy: 0.7432 - val_loss: 0.7991 - val_categorical_accuracy:
0.6719
Epoch 18/25
15/14 [=====] - 11s 714ms/step - loss: 0.6561 -
categorical_accuracy: 0.7365 - val_loss: 0.7956 - val_categorical_accuracy:
0.6660
Epoch 19/25

```

15/14 [=====] - 10s 697ms/step - loss: 0.6389 -
categorical_accuracy: 0.7427 - val_loss: 0.8070 - val_categorical_accuracy:
0.6641
Epoch 20/25
15/14 [=====] - 10s 697ms/step - loss: 0.6126 -
categorical_accuracy: 0.7609 - val_loss: 0.9036 - val_categorical_accuracy:
0.6777
Epoch 21/25
15/14 [=====] - 11s 705ms/step - loss: 0.6242 -
categorical_accuracy: 0.7484 - val_loss: 0.9776 - val_categorical_accuracy:
0.6602
Epoch 22/25
15/14 [=====] - 11s 703ms/step - loss: 0.5987 -
categorical_accuracy: 0.7661 - val_loss: 0.8078 - val_categorical_accuracy:
0.6621
Epoch 23/25
15/14 [=====] - 11s 701ms/step - loss: 0.5807 -
categorical_accuracy: 0.7635 - val_loss: 0.8075 - val_categorical_accuracy:
0.6621
Epoch 24/25
15/14 [=====] - 10s 694ms/step - loss: 0.5899 -
categorical_accuracy: 0.7708 - val_loss: 0.7806 - val_categorical_accuracy:
0.6562
Epoch 25/25
15/14 [=====] - 11s 704ms/step - loss: 0.5587 -
categorical_accuracy: 0.7703 - val_loss: 0.8241 - val_categorical_accuracy:
0.6582

```

7.3 Evaluating on test set

- show the accuracy

```
[ ]: ##YOUR CODE HERE
```

```
[27]: loss, accuracy = model2.evaluate([test_x, test_aspect], test_y)

print("Categorical accuracy: {:.3f}%".format(accuracy*100))
```

```

638/638 [=====] - 2s 3ms/step
Categorical accuracy: 59.718%

```

```
[ ]:
```